

Lecture 10: Classic Games

David Silver

Outline

- 1 State of the Art
- 2 Game Theory
- 3 Minimax Search
- 4 Self-Play Reinforcement Learning
- 5 Combining Reinforcement Learning and Minimax Search
- 6 Reinforcement Learning in Imperfect-Information Games
- 7 Conclusions

Why Study Classic Games?

- Simple rules, deep concepts
- Studied for hundreds or thousands of years
- Meaningful IQ test
- *Drosophila* of artificial intelligence
- Microcosms encapsulating real world issues
- Games are fun!

AI in Games: State of the Art

Program	Level of Play	Program to Achieve Level
Checkers	Perfect	<i>Chinook</i>
Chess	Superhuman	<i>Deep Blue</i>
Othello	Superhuman	<i>Logistello</i>
Backgammon	Superhuman	<i>TD-Gammon</i>
Scrabble	Superhuman	<i>Maven</i>
Go	Grandmaster	<i>MoGo</i> ¹ , <i>Crazy Stone</i> ² , <i>Zen</i> ³
Poker ⁴	Superhuman	<i>Polaris</i>

¹9 × 9

²9 × 9 and 19 × 19

³19 × 19

⁴Heads-up Limit Texas Hold'em

RL in Games: State of the Art

Program	Level of Play	RL Program to Achieve Level
Checkers	Perfect	<i>Chinook</i>
Chess	International Master	<i>KnightCap / Meep</i>
Othello	Superhuman	<i>Logistello</i>
Backgammon	Superhuman	<i>TD-Gammon</i>
Scrabble	Superhuman	<i>Maven</i>
Go	Grandmaster	<i>MoGo</i> ¹ , <i>Crazy Stone</i> ² , <i>Zen</i> ³
Poker ⁴	Superhuman	<i>SmooCT</i>

¹ 9×9

² 9×9 and 19×19

³ 19×19

⁴Heads-up Limit Texas Hold'em

Optimality in Games

- What is the optimal policy π^i for i th player?
- If all other players fix their policies π^{-i}
- **Best response** $\pi_*^i(\pi^{-i})$ is optimal policy against those policies
- **Nash equilibrium** is a joint policy for all players

$$\pi^i = \pi_*^i(\pi^{-i})$$

- such that every player's policy is a best response
- i.e. no player would choose to deviate from Nash

Single-Agent and Self-Play Reinforcement Learning

- Best response is solution to single-agent RL problem
 - Other players become part of the environment
 - Game is reduced to an MDP
 - Best response is optimal policy for this MDP
- Nash equilibrium is fixed-point of self-play RL
 - Experience is generated by playing games between agents

$$a_1 \sim \pi^1, a_2 \sim \pi^2, \dots$$

- Each agent learns best response to other players
- One player's policy determines another player's environment
- All players are adapting to each other

Two-Player Zero-Sum Games

We will focus on a special class of games:

- A **two-player game** has two (alternating) players
 - We will name player 1 *white* and player 2 *black*
- A **zero sum game** has equal and opposite rewards for black and white

$$R^1 + R^2 = 0$$

We consider methods for finding Nash equilibria in these games

- Game tree search (i.e. planning)
- Self-play reinforcement learning

Perfect and Imperfect Information Games

- A **perfect information** or **Markov** game is fully observed
 - Chess
 - Checkers
 - Othello
 - Backgammon
 - Go
- An **imperfect information** game is partially observed
 - Scrabble
 - Poker
- We focus first on perfect information games

Minimax

- A **value function** defines the expected total reward given joint policies $\pi = \langle \pi^1, \pi^2 \rangle$

$$v_{\pi}(s) = \mathbb{E}_{\pi} [G_t \mid S_t = s]$$

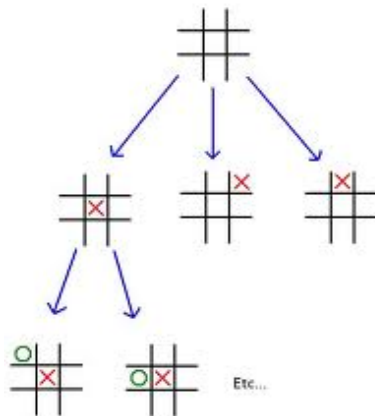
- A **minimax** value function maximizes white's expected return while minimizing black's expected return

$$v_*(s) = \max_{\pi^1} \min_{\pi^2} v_{\pi}(s)$$

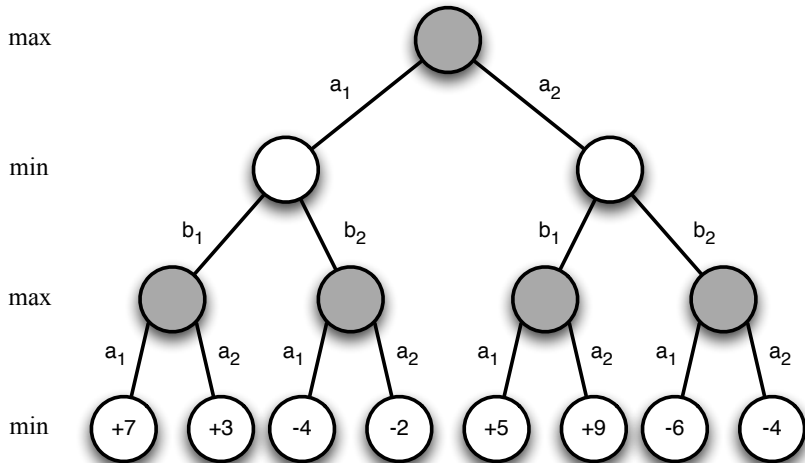
- A **minimax** policy is a joint policy $\pi = \langle \pi^1, \pi^2 \rangle$ that achieves the minimax values
- There is a unique minimax value function
- A minimax policy is a Nash equilibrium

Minimax Search

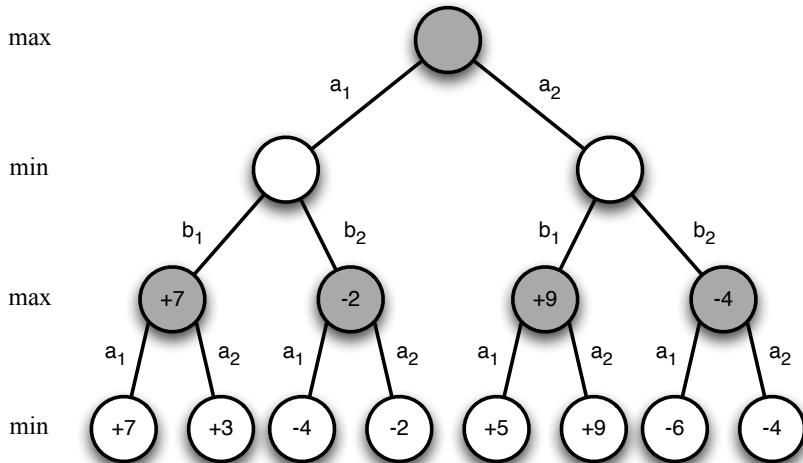
- Minimax values can be found by depth-first game-tree search
- Introduced by Claude Shannon: *Programming a Computer for Playing Chess*
- Ran on paper!



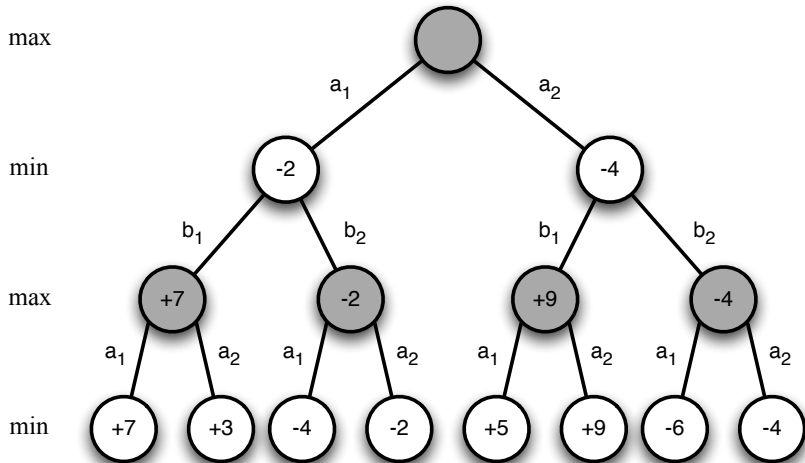
Minimax Search Example



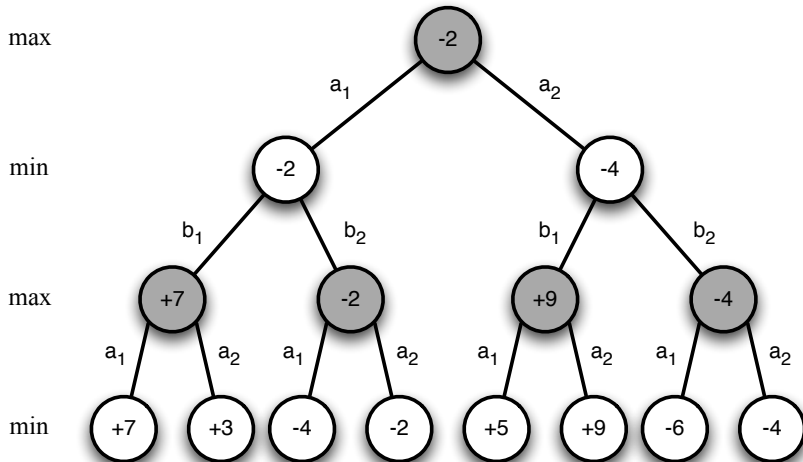
Minimax Search Example



Minimax Search Example



Minimax Search Example

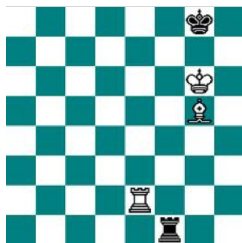



Value Function in Minimax Search

- Search tree grows exponentially
- Impractical to search to the end of the game
- Instead use value function approximator $v(s, \mathbf{w}) \approx v_*(s)$
 - aka *evaluation function*, *heuristic function*
- Use value function to estimate minimax value at leaf nodes
- Minimax search run to fixed depth with respect to leaf values

Binary-Linear Value Function

- Binary feature vector $\mathbf{x}(s)$: e.g. one feature per piece
- Weight vector \mathbf{w} : e.g. value of each piece
- Position is evaluated by summing weights of active features



$$v(s, \mathbf{w}) = \mathbf{x}(s) \cdot \mathbf{w} = \begin{bmatrix} 1 \\ 1 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ \vdots \end{bmatrix} \cdot \begin{bmatrix} +5 \\ +3 \\ +1 \\ -5 \\ -3 \\ -1 \\ \vdots \end{bmatrix}$$


$$v(s, \mathbf{w}) = 5 + 3 - 5 = 3$$

Deep Blue

- Knowledge
 - 8000 handcrafted chess features
 - Binary-linear value function
 - Weights largely hand-tuned by human experts
- Search
 - High performance parallel alpha-beta search
 - 480 special-purpose VLSI chess processors
 - Searched 200 million positions/second
 - Looked ahead 16-40 ply
- Results
 - Defeated human champion Garry Kasparov 4-2 (1997)
 - Most watched event in internet history

Chinook

- Knowledge
 - Binary-linear value function
 - 21 knowledge-based features (position, mobility, ...)
 - x4 phases of the game
- Search
 - High performance alpha-beta search
 - Retrograde analysis
 - Search backward from won positions
 - Store all winning positions in lookup tables
 - Plays perfectly from last n checkers
- Results
 - Defeated Marion Tinsley in world championship 1994
 - won 2 games but Tinsley withdrew for health reasons
 - Chinook *solved* Checkers in 2007
 - perfect play against God

Self-Play Temporal-Difference Learning

- Apply value-based RL algorithms to games of self-play
- MC: update value function towards the return G_t

$$\Delta \mathbf{w} = \alpha(G_t - v(S_t, \mathbf{w})) \nabla_{\mathbf{w}} v(S_t, \mathbf{w})$$

- TD(0): update value function towards successor value $v(S_{t+1})$

$$\Delta \mathbf{w} = \alpha(v(S_{t+1}, \mathbf{w}) - v(S_t, \mathbf{w})) \nabla_{\mathbf{w}} v(S_t, \mathbf{w})$$

- TD(λ): update value function towards the λ -return G_t^λ

$$\Delta \mathbf{w} = \alpha(G_t^\lambda - v(S_t, \mathbf{w})) \nabla_{\mathbf{w}} v(S_t, \mathbf{w})$$

Policy Improvement with Afterstates

- For deterministic games it is sufficient to estimate $v_*(s)$
- This is because we can efficiently evaluate the **afterstate**

$$q_*(s, a) = v_*(succ(s, a))$$

- Rules of the game define the successor state $succ(s, a)$
- Actions are selected e.g. by min/maximising afterstate value

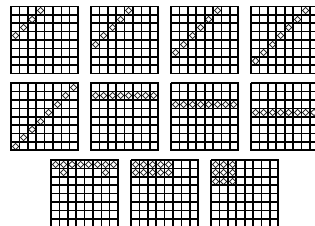
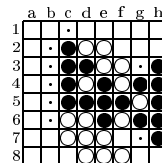
$$A_t = \operatorname{argmax}_a v_*(succ(S_t, a)) \quad \text{for white}$$

$$A_t = \operatorname{argmin}_a v_*(succ(S_t, a)) \quad \text{for black}$$

- This improves joint policy for both players

Self-Play TD in Othello: *Logistello*

- Logistello created its own features
- Start with raw input features, e.g. “black stone at C1?”
- Construct new features by conjunction/disjunction
- Created 1.5 million features in different configurations
- Binary-linear value function using these features



Reinforcement Learning in Logistello

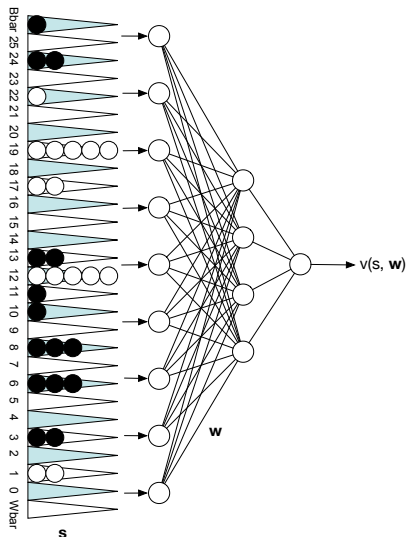
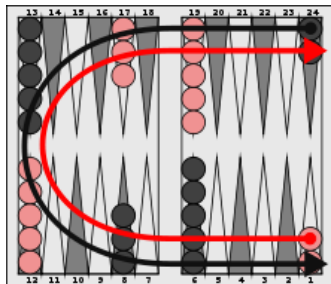
Logistello used generalised policy iteration

- Generate batch of self-play games from current policy
- Evaluate policies using Monte-Carlo (regress to outcomes)
- Greedy policy improvement to generate new players

Results

- Defeated World Champion Takeshi Murukami 6-0

TD Gammon: Non-Linear Value Function Approximation



Self-Play TD in Backgammon: *TD-Gammon*

- Initialised with random weights
- Trained by games of self-play
- Using non-linear temporal-difference learning

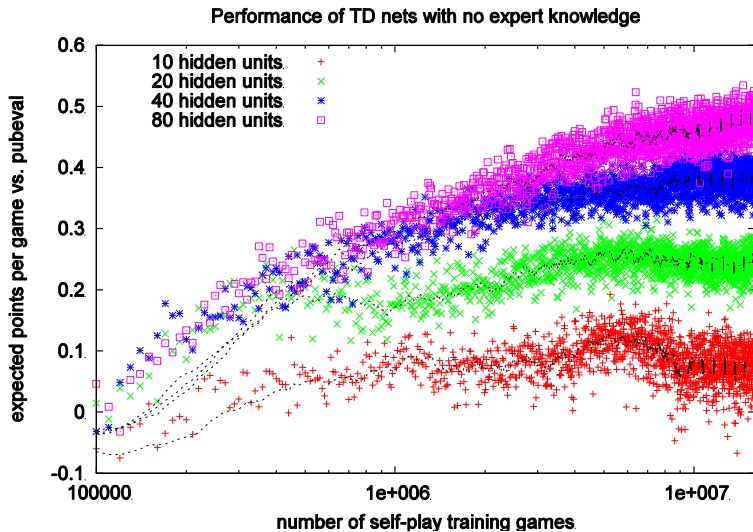
$$\delta_t = v(S_{t+1}, \mathbf{w}) - v(S_t, \mathbf{w})$$
$$\Delta \mathbf{w} = \alpha \delta_t \nabla_{\mathbf{w}} v(S_t, \mathbf{w})$$

- Greedy policy improvement (no exploration)
- Algorithm always converged in practice
- Not true for other games

TD Gammon: Results

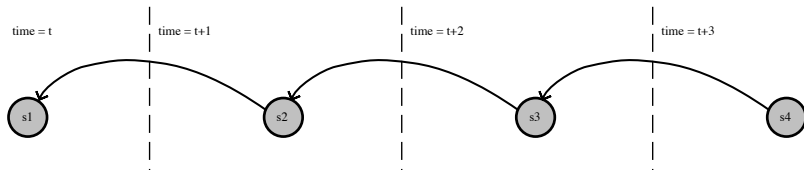
- Zero expert knowledge \implies strong intermediate play
- Hand-crafted features \implies advanced level of play (1991)
- 2-ply search \implies strong master play (1993)
- 3-ply search \implies superhuman play (1998)
- Defeated world champion Luigi Villa 7-1 (1992)

New TD-Gammon Results



Simple TD

- **TD**: update value towards successor value



- Value function approximator $v(s, \mathbf{w})$ with parameters \mathbf{w}
- Value function backed up from raw value at next state

$$v(S_t, \mathbf{w}) \leftarrow v(S_{t+1}, \mathbf{w})$$

- First learn value function by TD learning
- Then use value function in minimax search (no learning)

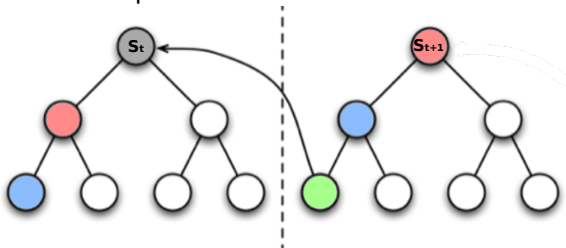
$$v_+(S_t, \mathbf{w}) = \min_{s \in \text{leaves}(S_t)} \max v(s, \mathbf{w})$$

Simple TD: Results

- Othello: superhuman performance in *Logistello*
- Backgammon: superhuman performance in *TD-Gammon*
- Chess: poor performance
- Checkers: poor performance
- In chess tactics seem necessary to find signal in position
- e.g. hard to find checkmates without search
- Can we learn directly from minimax search values?

TD Root

- **TD root:** update value towards successor search value



- Search value is computed at root position S_t

$$v_+(S_t, \mathbf{w}) = \min_{s \in \text{leaves}(S_t)} \max_{\mathbf{w}} v(s, \mathbf{w})$$

- Value function backed up from *search value* at next state

$$v(S_t, \mathbf{w}) \leftarrow v_+(S_{t+1}, \mathbf{w}) = v(l_+(S_{t+1}), \mathbf{w})$$

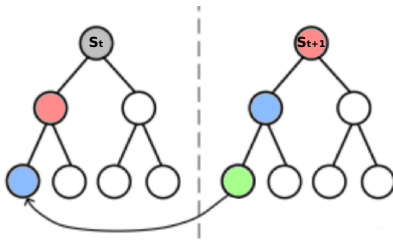
- Where $l_+(s)$ is the leaf node achieving minimax value from s

TD Root in Checkers: *Samuel's Player*

- First ever TD learning algorithm (*Samuel 1959*)
- Applied to a Checkers program that learned by self-play
- Defeated an amateur human player
- Also used other ideas we might now consider strange

TD Leaf

- **TD leaf**: update search value towards successor search value



- Search value computed at current and next step

$$v_+(S_t, \mathbf{w}) = \min_{s \in \text{leaves}(S_t)} v(s, \mathbf{w}), \quad v_+(S_{t+1}, \mathbf{w}) = \min_{s \in \text{leaves}(S_{t+1})} v(s, \mathbf{w})$$

- Search value at step t backed up from *search value* at $t + 1$

$$\begin{aligned} v_+(S_t, \mathbf{w}) &\leftarrow v_+(S_{t+1}, \mathbf{w}) \\ \implies v(l_+(S_t), \mathbf{w}) &\leftarrow v(l_+(S_{t+1}), \mathbf{w}) \end{aligned}$$

TD leaf in Chess: *Knightcap*

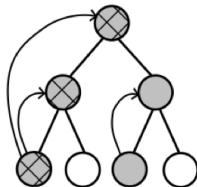
- Learning
 - *Knightcap* trained against expert opponent
 - Starting from standard piece values only
 - Learnt weights using TD leaf
- Search
 - Alpha-beta search with standard enhancements
- Results
 - Achieved master level play after a small number of games
 - Was not effective in self-play
 - Was not effective without starting from good weights

TD leaf in Checkers: *Chinook*

- Original Chinook used hand-tuned weights
- Later version was trained by self-play
- Using TD leaf to adjust weights
 - Except material weights which were kept fixed
- Self-play weights performed \geq hand-tuned weights
- i.e. learning to play at superhuman level

TreeStrap

- **TreeStrap**: update search values towards deeper search values



- Minimax search value computed at *all* nodes $s \in \text{nodes}(S_t)$
- Value backed up from search value, at same step, for all nodes

$$v(s, \mathbf{w}) \leftarrow v_+(s, \mathbf{w})$$

$$\implies v(s, \mathbf{w}) \leftarrow v(l_+(s), \mathbf{w})$$

Treestrap in Chess: *Meep*

- Binary linear value function with 2000 features
- Starting from random initial weights (no prior knowledge)
- Weights adjusted by TreeStrap
- Won 13/15 vs. international masters
- Effective in self-play
- Effective from random initial weights

Simulation-Based Search

- Self-play reinforcement learning can replace search
- Simulate games of self-play from root state S_t
- Apply RL to simulated experience
 - Monte-Carlo Control \implies Monte-Carlo Tree Search
 - Most effective variant is UCT algorithm
 - Balance exploration/exploitation in each node using UCB
 - Self-play UCT converges on minimax values
 - Perfect information, zero-sum, 2-player games
 - Imperfect information: see next section

Performance of MCTS in Games

- MCTS is best performing method in many challenging games
 - Go (last lecture)
 - Hex
 - Lines of Action
 - Amazons
- In many games simple Monte-Carlo search is enough
 - Scrabble
 - Backgammon

Simple Monte-Carlo Search in Maven

■ Learning

- Maven evaluates moves by $score + v(rack)$
- Binary-linear value function of rack
- Using one, two and three letter features
- Q??????, QU?????, III????
- Learnt by Monte-Carlo policy iteration (cf. Logistello)

■ Search

- Roll-out moves by imagining n steps of self-play
- Evaluate resulting position by $score + v(rack)$
- Score move by average evaluation in rollouts
- Select and play highest scoring move
- Specialised endgame search using B^*

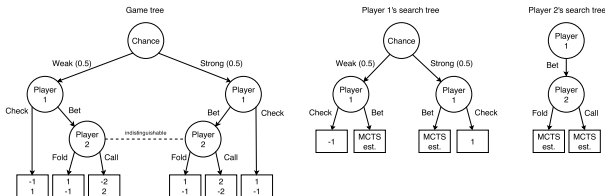
Maven: Results

- Maven beat world champion Adam Logan 9-5
- Here Maven predicted endgame to finish with MOUTHPART
- Analysis showed Maven had error rate of 3 points per game

M ₃	O ₁	U ₁	T ₁	H ₄		A ₁	R ₁	T ₁			2L			3W
A ₁	E ₁				3L				Q ₁₀				2W	
T ₁		2W				2L		2L	U ₁			G ₂		
H ₄	U ₁	R ₁	T ₁				2L		A ₁		2W	R ₁		2L
	N ₁	E ₁	O ₁	N ₁					I ₁	S ₁		E ₁		L ₁
	3L		D ₂	O ₁	Z ₁₀	Y ₄			3L	P ₃		A ₁	X ₈	E ₁
		E ₁				E ₁		2L	J ₈	A ₁	H ₄	S ₁		I ₁
I ₁	A ₁	M ₃	B ₃		C ₃	A ₁	V ₄	Y ₄		N ₁	2L	E ₁		3W
	H ₄	E ₁				R ₁		2L		K ₅		2L		
	3L	N ₁		F ₄	3L	L ₁			B ₃				3L	
		D ₂		E ₁		O ₁			O ₁	R ₁				
2L	D ₂	E ₁	V ₄	I ₁	A ₁	N ₁	C ₃	E ₁	S ₁		2W			2L
		D ₂		G ₂		G ₂	O ₁	2L				2W		
	2W			N ₁	3L		F ₄		3L				2W	
P ₃	I ₁	L ₁	I ₁	S ₁			T ₁	U ₁	T ₁	O ₁	R ₁	I ₁	A ₁	L ₁

Game-Tree Search in Imperfect Information Games

- Players have different information states and therefore separate search trees



- There is one node for each information state
 - summarising what a player knows
 - e.g. the cards they have seen
- Many real states may share the same information state
- May also aggregate states e.g. with similar value

Solution Methods for Imperfect Information Games

Information-state game tree may be solved by:

- Iterative forward-search methods
 - e.g. Counterfactual regret minimization
 - “Perfect” play in Poker (heads-up limit Hold'em)
- Self-play reinforcement learning
- e.g. Smooth UCT
 - 3 silver medals in two- and three-player Poker (limit Hold'em)
 - Outperformed massive-scale forward-search agents

Smooth UCT Search

- Apply MCTS to information-state game tree
- Variant of UCT, inspired by game-theoretic Fictitious Play
 - Agents learn against and respond to opponents' average behaviour
- Extract average strategy from nodes' action counts,
 $\pi_{avg}(a|s) = \frac{N(s,a)}{N(s)}.$
- At each node, pick actions according to

$$A \sim \begin{cases} \text{UCT}(S), & \text{with probability } \eta \\ \pi_{avg}(\cdot|S), & \text{with probability } 1 - \eta \end{cases}$$

- Empirically, in variants of Poker:
 - Naive MCTS diverged
 - Smooth UCT converged to Nash equilibrium

RL in Games: A Successful Recipe

Program	Input features	Value Fn	RL	Training	Search
Chess <i>Meep</i>	Binary <i>Pieces, pawns, ...</i>	Linear	TreeStrap	Self-Play / Expert	$\alpha\beta$
Checkers <i>Chinook</i>	Binary <i>Pieces, ...</i>	Linear	TD leaf	Self-Play	$\alpha\beta$
Othello <i>Logistello</i>	Binary <i>Disc configs</i>	Linear	MC	Self-Play	$\alpha\beta$
Backgammon <i>TD Gammon</i>	Binary <i>Num checkers</i>	Neural network	TD(λ)	Self-Play	$\alpha\beta$ / MC
Go <i>MoGo</i>	Binary <i>Stone patterns</i>	Linear	TD	Self-Play	MCTS
Scrabble <i>Maven</i>	Binary <i>Letters on rack</i>	Linear	MC	Self-Play	MC search
Limit Hold'em <i>SmooCT</i>	Binary <i>Card abstraction</i>	Linear	MCTS	Self-Play	-